

Group 5

Code for this assignment at
<https://github.com/TetroVolt/COSC3337Assign2>

Who we are




Raymond Sutrisno



Niels Moeller




Gal Egozi



Colby Kuhnel

Data Set I (Molecular Biology Data Set)

Data Set Characteristics:	Sequential, Domain-Theory	Number of Instances:	3190	Area:	Life
Attribute Characteristics:	Categorical	Number of Attributes:	61	Date Donated	1992-01-01
Associated Tasks:	Classification	Missing Values?	No	Number of Web Hits:	95406



- Introns do not code for proteins
- Exons code for proteins

Data Set II: Preliminary analysis

1. Each Sequence is 60 base pairs long
2. 8 possible base pairs: ('A', 'C', 'D', 'G', 'N', 'R', 'S', 'T')
3. Three Class Labels
 - a. Exon-Intron Boundaries (EI) (767 examples)
 - b. Intron-Exon Boundaries (IE) (768 examples)
 - c. Neither (N) (1655 examples)

Data Preprocessing I (RAW data)

A1    Class

	A	B	C
1	Class	Source	Sequence
2	EI	ATRINS-DONOR-521	CCAGCTGCATCACAGGAGGCCAGCGAGCAGGTCTGTTCCAAGGGCCTTCGAGCCAGTCTG
3	EI	ATRINS-DONOR-905	AGACCCGCCGGGAGGCGGAGGACCTGCAGGGTGAGCCCCACCGCCCCTCCGTGCCCCCGC
4	EI	BABAPOE-DONOR-30	GAGGTGAAGGACGTCTTCCCCAGGAGCCGGTGAGAAGCGCAGTCGGGGGCACGGGGATG
5	EI	BABAPOE-DONOR-867	GGGCTGCGTTGCTGGTCACATTCTGGCAGGTATGGGGCGGGGCTTGCTCGGTTTTCCCC
6	EI	BABAPOE-DONOR-2817	GCTCAGCCCCCAGGTACCCAGGAACTGACGTGAGTGTCCTCCATCCCGGCCCTTGACCCT
7	EI	CHPIGECA-DONOR-378	CAGACTGGGTGGACAACAAAACCTTCAGCGGTAAGAGAGGGCCAAGCTCAGAGACCACAG
8	EI	CHPIGECA-DONOR-903	CCTTTGAGGACAGCACCAAGAAGTGTGCAGGTACGTTCCACCTGCCCTGGTGCCGCCA
9	EI	CHPIGECA-DONOR-1313	CCCTCGTGCGGTCCACGACCAAGACCAGCGGTGAGCCACGGGCAGGCCGGGGTCGTGGGG
10	EI	GCRHBBA1-DONOR-1260	TGGCGACTACGGCGCGGAGGCCCTGGAGAGGTGAGGACCCTCCTGTCCCTGCTCCAGTCC
11	EI	GCRHBBA1-DONOR-1590	AAGCTGACAGTGGAACCCGGTCAACTTCAAGGTGAGCCAGGAGTCGGGTGGGAGGGTGAGA
12	EI	GCRHBBA6-DONOR-461	TGGCGACTACGGCGCGGAGGCCCTGGAGAGGTGAGGACCCTGGTATCCCTGCTGCCAGTC
13	EI	GCRHBBA6-DONOR-795	AAGCTGAGAGTGGAACCTGTCAACTTCAAGGTGAGCCACCAGTCGGGTGGGGAGGGTGAG
14	EI	GIBHBGGL-DONOR-2278	GGAAGATGCTGGAGGAGAAACCCTGGGAAGGTAGGCTCTGGTGACCAGGACAAGGGAGGG
15	EI	GIBHBGGL-DONOR-2624	AAGCTGCATGTGGATCCTGAGAACTTCAGGGTGAGTACAGGAGATGTTTCAGCCCTGTTG
16	EI	GIBHBGGL-DONOR-7198	GGAAGATGTTGGAGGAGAAACCCTGGGAAGGTAGGCTCTGGTGACCAGGACAAGGGAGGG
17	EI	GIBHBGGL-DONOR-7544	AAGCTGCATGTGGATCCTGAGAACTTCAGGGTGAGTACAGGAGATGTTTCAGCCCTGTTG
18	EI	HIIM414TP-DONOR-1072	GGCACCACCACTGACCTGGGACAGTGAAATCGTAAGTATGCTTTCACTGCGAGGGGTTCT

Algorithms considered

Required:

1. Neural Networks (Gal Egozi)
2. Support Vector Machines (Colby Kuhnel)

Chosen freely:

1. Decision Tree (Niels Moeller)
2. Random Forest (Raymond Sutrisno)

Data Processing II (considerations)

1. NN and SVM are “One vs. All” algorithms
2. Features are categorical in nature
3. Need One Hot Encoding for both class labels and features

Data Processing III (procedure)

1. One hot encode features
 - a. Each Sequence is 60 base pairs long and categorical
 - b. 60 features
 - c. 8 base pair categories ('A', 'C', 'D', 'G', 'N', 'R', 'S', 'T')
 - d. $60 * 8 = 480$ columns total after one hot encoding
2. One hot encode classes for algorithms that need it
 - a. Three classes: 'EI' , 'IE', 'N'

Training Procedure

1. Gridsearch model parameters to tune for the best model using stratified 10 fold cross validation
2. Report accuracies (CV mainly, training and test)
3. Optional metrics considered
 - a. Confusion Matrices
 - b. Learning Curves

Implementation Info

1. Language used: Python
2. Utility Libraries:
 - a. Pandas (data processing)
 - b. Numpy (data processing)
 - c. Scikit Learn (data processing)
3. Machine Learning Libraries
 - a. Scikit Learn (SVM, Random Forest, Decision Tree)
 - b. Keras (Neural networks)

Results

Random Forest (trees based on C4.5)

Random Forest I (model parameter space available)

1. Criterion ['gini', 'entropy']
2. Max_depth
3. Min sample splits
4. Min impurity decrease
5. ...

Random Forest I (model parameter space used)

1. `'N_estimators': [25, 50, 100, 150]`
2. `'max_depth': [2, 4, 8, 16, 32, 64]`

Random Forest II (Raw Output) (20 % Test)

Dataset characteristics:

Number of examples in the dataset = 3190

Number of examples reserved for test set = 638

Number of examples reserved for training via 10 fold CV = 2552

Class Distribution Ratio (N : EI : IE) = 2 : 1 : 1

N features : 60, all categorical (DNA base pairs in 60 base pair long sequence)

Grid Search parameter space for Random Forest =

`{'n_estimators': [25, 50, 100, 150], 'max_depth': [2, 4, 8, 16, 32, 64]}`

best estimator parameters found = `{'max_depth': 32, 'n_estimators': 150}`

best estimator mean training score = 0.9995646116203684

best estimator mean validation score = 0.9114420062695925

best estimator test score = 0.9247648902821317

Random Forest III (Raw Output) (50 % Test)

Dataset characteristics:

Number of examples in the dataset = 3190

Number of examples reserved for test set = 1595

Number of examples reserved for training via 10 fold CV = 1595

Class Distribution Ratio (N : EI : IE) = 2 : 1 : 1

N features : 60, all categorical (DNA base pairs in 60 base pair long sequence)

Grid Search parameter space for Random Forest =

`{'n_estimators': [25, 50, 100, 150], 'max_depth': [2, 4, 8, 16, 32, 64]}`

best estimator parameters found = `{'max_depth': 16, 'n_estimators': 100}`

best estimator mean training score = 0.9992337891743421

best estimator mean validation score = 0.8915360501567398

best estimator test score = 0.8984326018808777

Random Forest IV (Raw Output) (70 % Test)

Dataset characteristics:

Number of examples in the dataset = 3190

Number of examples reserved for test set = 2233

Number of examples reserved for training via 10 fold CV = 957

Class Distribution Ratio (N : EI : IE) = 2 : 1 : 1

N features : 60, all categorical (DNA base pairs in 60 base pair long sequence)

Grid Search parameter space for Random Forest =

`{'n_estimators': [25, 50, 100, 150], 'max_depth': [2, 4, 8, 16, 32, 64]}`

best estimator parameters found = `{'max_depth': 16, 'n_estimators': 100}`

best estimator mean training score = 1.0

best estimator mean validation score = 0.8610240334378265

best estimator test score = 0.8669950738916257

Random Forest Best Parameters

Train : Test Ratio	Max Depth	N Estimators
80:20	32	150 (MAX)
50:50	16	100
30:70	16	100

Random Forest Mean Scores

Train : Test Ratio	Mean Train	Mean CV	Test Score
80:20	0.9995	0.9114	0.9247
50:50	0.9992	0.8915	0.8984
30:70	1.0000	0.8610	0.8669

Random Forest Mean Scores

Train : Test Ratio	Mean Train	Mean CV	Test Score	Max Depth	N Estimators
80:20	0.9995	0.9114	0.9247	32	150
50:50	0.9992	0.8915	0.8984	16	100
30:70	1.0000	0.8610	0.8669	16	100

Decision Tree (SKLearn based on C4.5)

Decision Trees 80/20

Dataset characteristics:

Number of examples in the dataset = 3190

Number of examples reserved for test set = 638

Number of examples reserved for training via 10 fold CV = 2552

Class Distribution Ratio (N : EI : IE) = 2 : 1 : 1

N features : 60, all categorical (DNA base pairs in 60 base pair long sequence)

Grid Search parameter space for Random Forest = {'max_depth': range(3, 20)}

best estimator parameters found = {'max_depth': 6}

best estimator mean training score = 0.9573321284086651

best estimator mean validation score = 0.9455329153605015

best estimator test score = 0.9404388714733543



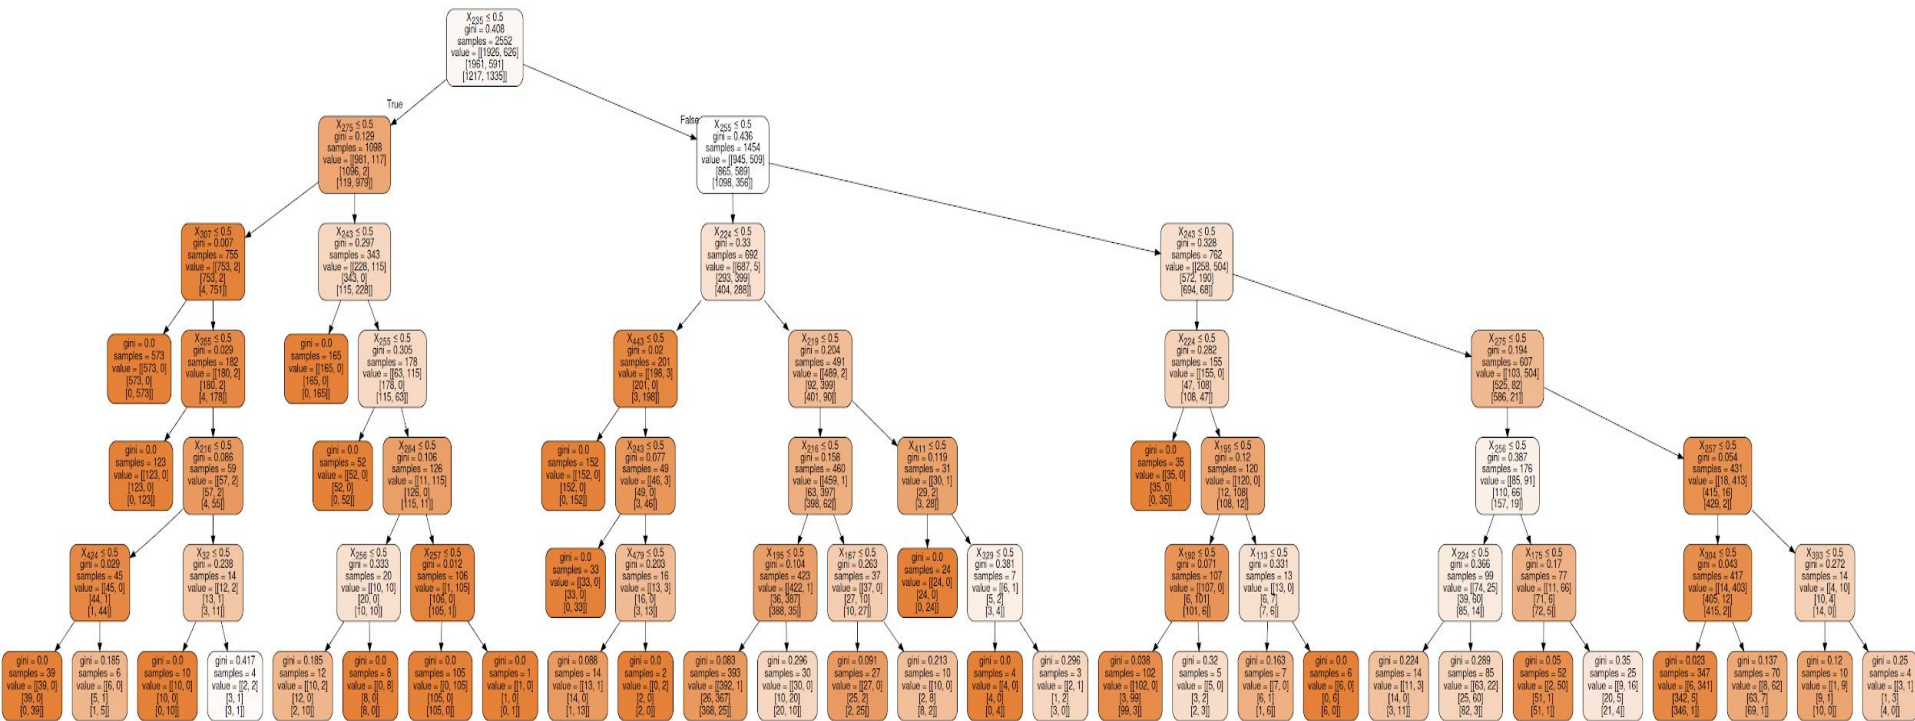
Decision Tree Scores

Train : Test Ratio	Mean Train	Mean CV	Test Score
80:20	0.9570	0.9447	0.9341
50:50	0.9484	0.9347	0.9373
30:70	0.9564	0.9362	0.9305

Decision Tree Best Parameters

Train : Test Ratio	Max Depth (3-20)	Min samples split [3,5,25]
80:20	5	25
50:50	5	3
30:70	6	3

Resulting Decision Tree



SVM

Support Vector Machines (20% Test)

Grid Search parameter space for SVM = [{'C': [100, 1000, 10000], 'gamma': [0.01, 0.001], 'kernel': ['poly', 'rbf'], 'degree': [2, 3, 4]}]

best estimator parameters found = {'C': 100, 'degree': 3, 'gamma': 0.01, 'kernel': 'poly'}

best estimator mean training score = 0.9996081466330754

best estimator mean validation score = 0.9678683385579937

best estimator test score = 0.9702194357366771

Process Time = 8.436292 Minutes

Support Vector Machines (50% Test)

Grid Search parameter space for SVM = [{'C': [100, 1000, 10000], 'gamma': [0.01, 0.001], 'kernel': ['poly', 'rbf'], 'degree': [2, 3, 4]}]

best estimator parameters found = {'C': 100, 'degree': 3, 'gamma': 0.01, 'kernel': 'poly'}

best estimator mean training score = 0.9994427513515088

best estimator mean validation score = 0.9605015673981191

best estimator test score = 0.9661442006269593

Process Time = 3.845329 Minutes

Support Vector Machines (70% Test)

Grid Search parameter space for SVM = [{'C': [100, 1000, 10000], 'gamma': [0.01, 0.001], 'kernel': ['poly', 'rbf'], 'degree': [2, 3, 4]}]

best estimator parameters found = {'C': 100, 'degree': 3, 'gamma': 0.01, 'kernel': 'poly'}

best estimator mean training score = 1.0

best estimator mean validation score = 0.9540229885057471

best estimator test score = 0.961486789072996

Process Time = 1.603256 Minutes

Support Vector Machines Mean Scores

Train : Test Ratio	Mean Train	Mean CV	Test Score
80:20	0.9996	0.9679	0.9702
50:50	0.9994	0.9605	0.9661
30:70	1.0	0.9540	0.9615






Neural Network

Neural Network







- One layer, 50 nodes
- Scores $95.52 \pm 0.74\%$ on stratified cross validation
- 20% test
 - 99.61% Train
 - 94.36% Test
- 50% test
 - 99.44% Train
 - 93.67 Test
- 70% test
 - 100% train
 - 91.49% test

Performance Comparison




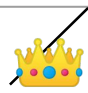


Top Dog for 20% Test

Algorithm	Best Test Score	Accompanying CV Score	Accompanying Train Score
Random Forest	0.9247	0.9114	0.9995
Decision Tree	0.9373	0.9347	0.9484
SVM   100	0.9702 	0.9679 	0.9996 
Neural Network	0.9436	0.9552	0.9961

Top Dog for 50% Test

Algorithm	Best Test Score	Accompanying CV Score	Accompanying Train Score
Random Forest	0.8984	0.8915	0.9992
Decision Tree	0.9373	0.9347	0.9484
SVM   	0.9661 	0.9605 	0.9994 
Neural Network	0.9436	0.9552	0.9961

Top Dog for 70% Test

Algorithm	Best Test Score	Accompanying CV Score	Accompanying Train Score
Random Forest	0.8669	0.8610	1.0000 
Decision Tree	0.9373	0.9347	0.9484
SVM   <u>100</u>	0.9615 	0.9540	1.0000 
Neural Network	0.9436	0.9552 	0.9961

What could have been done differently?

- Investigated Gini importances of features
- Investigated other pattern recognition neural network architectures
- Investigated learning curves of models to better understand generalization behaviors
- Investigated Confusion matrices to see if class distribution was an issue that skewed different types of errors
- Used stratified K fold rather than random sample K fold
- Played more with parameters in general

Thank you!

End of presentation